

# End to End Hospital Analytics

## Overview

This project delivers a comprehensive hospital analytics project developed using SQL for data processing and querying, and Power BI for data modeling and visualization. Designed to reflect the complexities of real world hospital operations, it provides insights into patient flow, physician performance, financial trends, and operational challenges through effective data storytelling.

With a background in healthcare, I approached this project with a focus on business relevance and user centric dashboard design. The outcome is a practical and actionable analytics resource that enables hospital management to track key performance indicators, identify inefficiencies, and support data driven decision making.

## Thought Process & Project Objectives

My starting point was to ask: *“What would a hospital administrator or stakeholder need to see daily to improve performance?”* From there, I outlined the following objectives:

- Consolidate diverse datasets covering clinical, operational, and financial data
- Validate data quality and ensure referential integrity across all available data
- Use SQL to uncover trends in patient behavior, treatment outcomes, doctor productivity, and billing health
- Design a Power BI dashboard that is easy to interpret, with distinct pages for clinical and financial insights, making it actionable for hospital stake holders
- Use storytelling to transform numbers into insight and strategy

## Data Sources

The analysis was based on five interconnected datasets:

- **patients:** basic demographics, insurance information, registration date
- **doctors:** specialization, experience, branch
- **appointments:** dates, reasons for visit, assigned doctor, attendance status
- **treatments:** type, cost, description, and date
- **billing:** amount, payment method, payment status

## Tools & Technologies

- **SQL (MySQL):** data cleaning, integrity checks, KPI computation, complex joins
- **Power BI:** dashboard building and DAX for aggregations
- **Excel:** initial exploration and sample mock-ups

# Data Cleaning and Validation

To ensure the accuracy and reliability of the hospital dataset, a thorough data cleaning and validation process was performed across all key tables: patients, doctors, appointments, treatments, and billing. The following validation steps were implemented:

## Missing Values Check

Each table was scanned for rows with missing critical fields using a combination of CONCAT\_WS() and IS NULL. This allowed identification of any records with incomplete information across key attributes such as patient details, appointment metadata, or billing amounts.

```
1 #1. Checking for columns with missing values across all datasets
2 -- This query checks for NULLs across all important columns in each table
3 SELECT 'patients' AS table_name, COUNT(*) AS null_rows
4 FROM patients
5 WHERE CONCAT_WS('', patient_id, first_name, last_name, gender, date_of_birth,
6                 contact_number, address, registration_date,
7                 insurance_provider, insurance_number) IS NULL
8 UNION ALL
9 SELECT 'doctors', COUNT(*) FROM doctors
10 WHERE CONCAT_WS('', doctor_id, first_name, last_name, specialization,
11                 phone_number, years_experience, hospital_branch) IS NULL
12 UNION ALL
13 SELECT 'appointments', COUNT(*) FROM appointments
14 WHERE CONCAT_WS('', appointment_id, patient_id, doctor_id, appointment_date,
15                 appointment_time, reason_for_visit, status) IS NULL
16 UNION ALL
17 SELECT 'treatments', COUNT(*) FROM treatments
18 WHERE CONCAT_WS('', treatment_id, appointment_id, treatment_type, cost, treatment_date) IS NULL
19 UNION ALL
20 SELECT 'billing', COUNT(*) FROM billing
21 WHERE CONCAT_WS('', bill_id, patient_id, treatment_id, bill_date, amount,
22                 payment_method, payment_status) IS NULL;
```

table_name	null_rows	
abc Filter...	abc Filter...	
patients	0	
doctors	0	
appointments	0	
treatments	0	
billing	0	

Result: No missing values were found in any table.

## Duplicate Record Detection

A duplicate check was executed by grouping each table on all its columns and identifying any repeated rows using the `HAVING COUNT(*) > 1` clause. This helped ensure that no redundant records were stored in the system.

```
1 #2. Checking for full duplicate rows across all tables
2 SELECT 'patients' AS table_name, COUNT(*) AS duplicate_rows
3 FROM patients
4 GROUP BY patient_id, first_name, last_name, gender, date_of_birth,
5         contact_number, address, registration_date, insurance_provider, insurance_number, email
6 HAVING COUNT(*) > 1
7 UNION ALL
8 SELECT 'doctors', COUNT(*) FROM doctors
9 GROUP BY doctor_id, first_name, last_name, specialization, phone_number, years_experience, hospital_branch, email
10 HAVING COUNT(*) > 1
11 UNION ALL
12 SELECT 'appointments', COUNT(*) FROM appointments
13 GROUP BY appointment_id, patient_id, doctor_id, appointment_date, appointment_time, reason_for_visit, status
14 HAVING COUNT(*) > 1
15 UNION ALL
16 SELECT 'treatments', COUNT(*) FROM treatments
17 GROUP BY treatment_id, appointment_id, treatment_type, description, cost, treatment_date
18 HAVING COUNT(*) > 1
19 UNION ALL
20 SELECT 'billing', COUNT(*) FROM billing
21 GROUP BY bill_id, patient_id, treatment_id, bill_date, amount, payment_method, payment_status
22 HAVING COUNT(*) > 1;
```

table_name	duplicate_rows	
abc Filter...	abc Filter...	

*Result: All tables were free from full duplicate rows.*

## Foreign Key Validation

To maintain referential integrity, foreign key relationships were validated across tables using `LEFT JOIN` checks. This ensured, for example, that every appointment was linked to a valid patient and doctor, and each billing entry referenced a valid treatment and patient.

```
1  #3. Validating foreign key relationships
2  -- Checks that foreign keys reference existing records
3  SELECT 'appointments - invalid patient' AS issue, a.appointment_id
4  FROM appointments a
5  LEFT JOIN patients p ON a.patient_id = p.patient_id
6  WHERE p.patient_id IS NULL
7  UNION ALL
8  SELECT 'appointments - invalid doctor', a.appointment_id
9  FROM appointments a
10 LEFT JOIN doctors d ON a.doctor_id = d.doctor_id
11 WHERE d.doctor_id IS NULL
12 UNION ALL
13 SELECT 'treatments - invalid appointment', t.treatment_id
14 FROM treatments t
15 LEFT JOIN appointments a ON t.appointment_id = a.appointment_id
16 WHERE a.appointment_id IS NULL
17 UNION ALL
18 SELECT 'billing - invalid treatment', b.bill_id
19 FROM billing b
20 LEFT JOIN treatments t ON b.treatment_id = t.treatment_id
21 WHERE t.treatment_id IS NULL
22 UNION ALL
23 SELECT 'billing - invalid patient', b.bill_id
24 FROM billing b
25 LEFT JOIN patients p ON b.patient_id = p.patient_id
26 WHERE p.patient_id IS NULL;
```

issue	appointment_id	
<input type="text" value="a b c Filter..."/>	<input type="text" value="a b c Filter..."/>	
No data		

*Result: No foreign key mismatches were identified.*

## Billing Consistency Check

Billing records marked as 'Paid' with an amount of zero were flagged for review, as they could indicate data entry or transactional errors.



```
1 #4. Checking for billing records marked as 'Paid' but with amount = 0
2 SELECT * FROM billing
3 WHERE payment_status = 'Paid' AND amount = 0;
```

bill_id	patient_id	treatment_id	bill_date	amount	payment_method
Filter...	Filter...	Filter...	Filter...	Filter...	Filter...
No data					

*Result: No such inconsistencies were found.*

## Insights & KPIs

Using SQL, I answered key business questions:

- **Total Patients:** 50
- **Total Appointments:** 200
- **Average Appointments per Patient:** 4.17
- **Total Revenue (Paid):** 551,000
- **Average Revenue per Patient:** 11,500
- **Failed Payment Rate:** 33%

Other queries revealed:

### Patients who never returned after registration (retention issue)

```
1 #8. Patients with no appointments
2 -- Useful for outreach or retention strategies
3 SELECT p.patient_id, CONCAT(p.first_name, ' ', p.last_name) AS patient_name
4 FROM patients p
5 LEFT JOIN appointments a ON p.patient_id = a.patient_id
6 WHERE a.appointment_id IS NULL;
```

patient_id	patient_name	
<input type="text" value="a b c Filter..."/>	<input type="text" value="a b c Filter..."/>	
P006	Linda Jones	
P015	Sarah Johnson	

*This query was designed to identify registered patients who never scheduled an appointment. Highlighting these gaps supports targeted follow up efforts and helps the hospital improve patient engagement and retention.*

## Top patients who visited multiple doctors (fragmented care or complex needs)

```
1  -- identifying highly engaged or complex patients
2  SELECT
3      a.patient_id,
4      CONCAT(p.first_name, ' ', p.last_name) AS patient_name,
5      COUNT(DISTINCT a.doctor_id) AS doctors_count
6  FROM appointments a
7  JOIN patients p ON a.patient_id = p.patient_id
8  GROUP BY a.patient_id, p.first_name, p.last_name
9  ORDER BY doctors_count DESC;
10
```

patient_id	patient_name	doctors_count
<input type="text"/> Filter...	<input type="text"/> Filter...	<input type="text"/> Filter...
P012	Laura Davis	8
P049	David Moore	6
P026	John Taylor	5
P036	Michael Wilson	5
P029	David Smith	5
P005	David Wilson	5
P035	David Wilson	5
P023	Linda Johnson	5
P019	Sarah Miller	4
P033	Michael Wilson	4
P007	Alex Johnson	4

*This analysis identifies patients who have seen multiple doctors, which may indicate complex medical needs, fragmented care, or inconsistent follow up. For instance, Patient P012 (Laura Davis) consulted with eight different doctors; suggesting the need to review her care pathway. Recognizing such patterns helps improve care coordination and can support targeted case management strategies.*



Insurance provider dominance across treatments

```
1  #10. Insurance distribution among patients
2  SELECT
3      p.insurance_provider,
4      COUNT(t.treatment_id) AS treatment_count,
5      ROUND(
6          COUNT(t.treatment_id) * 100.0 /
7          SUM(COUNT(t.treatment_id)) OVER (),
8          2
9      ) AS percentage_share
10 FROM treatments t
11 JOIN appointments a ON t.appointment_id = a.appointment_id
12 JOIN patients p ON a.patient_id = p.patient_id
13 GROUP BY p.insurance_provider
14 ORDER BY treatment_count DESC;
15
```

insurance_provi...	treatment_count	percentage_share
abc Filter...	abc Filter...	abc Filter...
MedCare Plus	84	42.00
WellnessCorp	58	29.00
PulseSecure	36	18.00
HealthIndia	22	11.00

*This query analyzes treatment volume by insurance provider to understand payer impact on hospital services. MedCare Plus accounts for 42% of all treatments; highlighting its strategic importance for contract negotiations and service planning.*

Treatment cost variability and which treatments drove revenue

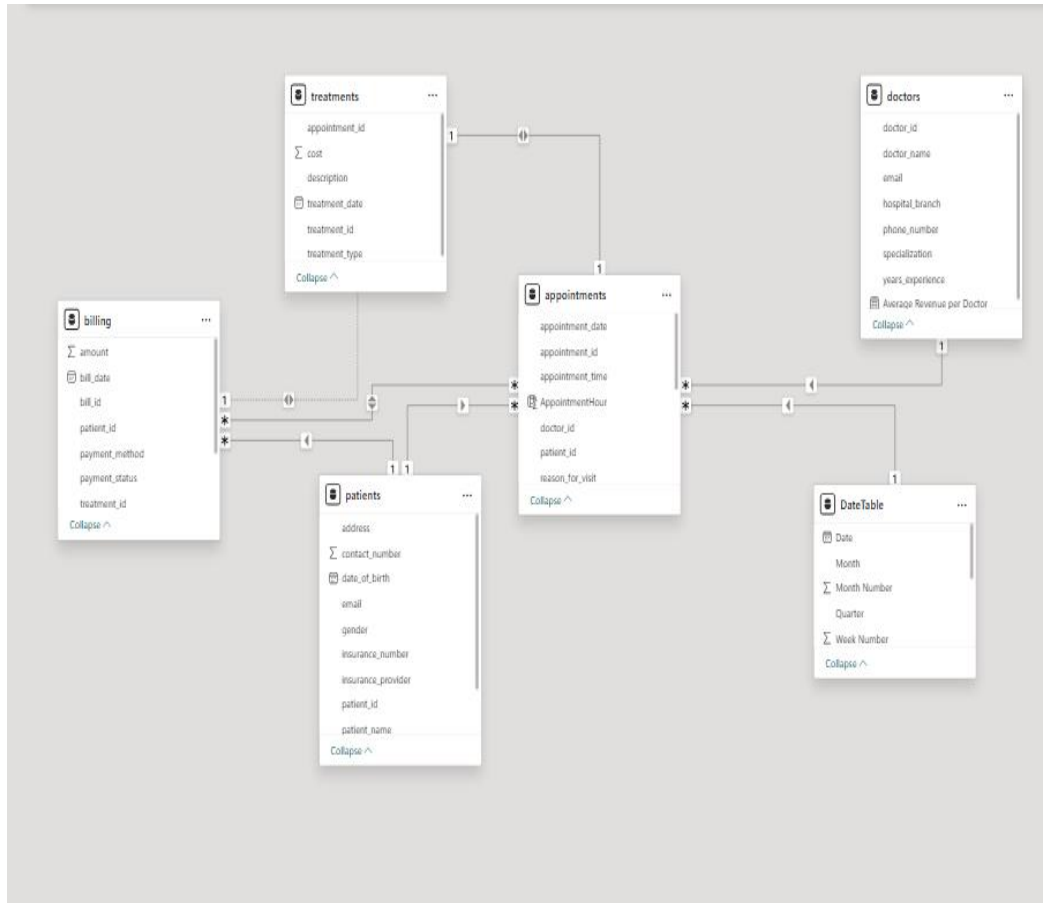
```
1  #15. Average treatment cost by treatment type
2  -- Shows pricing trends and potential over or undercharging
3  SELECT treatment_type, ROUND(AVG(cost), 2) AS avg_cost
4  FROM treatments
5  GROUP BY treatment_type
6  ORDER BY avg_cost DESC;
```

treatment_type	avg_cost	
abc Filter...	abc Filter...	
MRI	3224.95	
Physiotherapy	2761.61	
X-Ray	2698.87	
Chemotherapy	2629.71	
ECG	2532.22	

*This query calculates the average cost per treatment type to uncover pricing patterns. MRI and Physiotherapy are the most expensive on average, which may warrant a review of resource allocation, pricing strategy, or patient accessibility for these services.*

## Dashboard Structure in Power BI

After confirming that all tables were clean, I validated foreign key connections between them. I built a **star schema** model with `appointments` and `billing` as central fact tables supported by dimensional patient and doctor information.



The image above showcases the star schema used in the dashboard. Fact tables `appointments` connected to dimension tables `patients`, `doctors`, `billing` and `treatments`, enabling efficient filtering and scalable reporting across clinical and financial use cases.

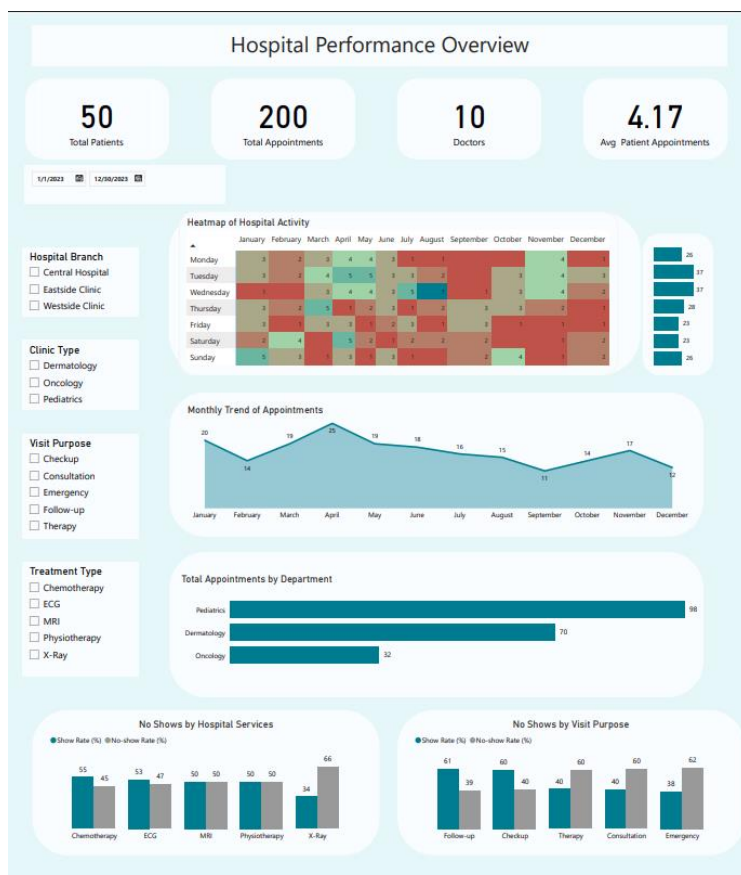
To ensure clarity and usability, the Power BI dashboard was split into two main pages:

## Page 1: Clinical Performance & Patient Flow

This page answers the question: “How are patients using our services and where are we losing efficiency?”

### Visuals:

- KPI Cards: total patients, average appointments, total appointments
- Heatmap: appointments by weekday (Tuesday and Wednesday were busiest)
- Line Graph: monthly appointment trend (peak in April, drop in September)
- Column Charts:
  - Appointments by specialization (Pediatrics had the highest volume)
  - Show vs No-Show rates by treatment type (X-ray 66% no-show)
  - Attendance by visit reason (Emergency and Therapy both ~60% no-show)
- Slicers: by hospital branch, doctor specialization, treatment type, date



### Hospital Overview and Operations Page

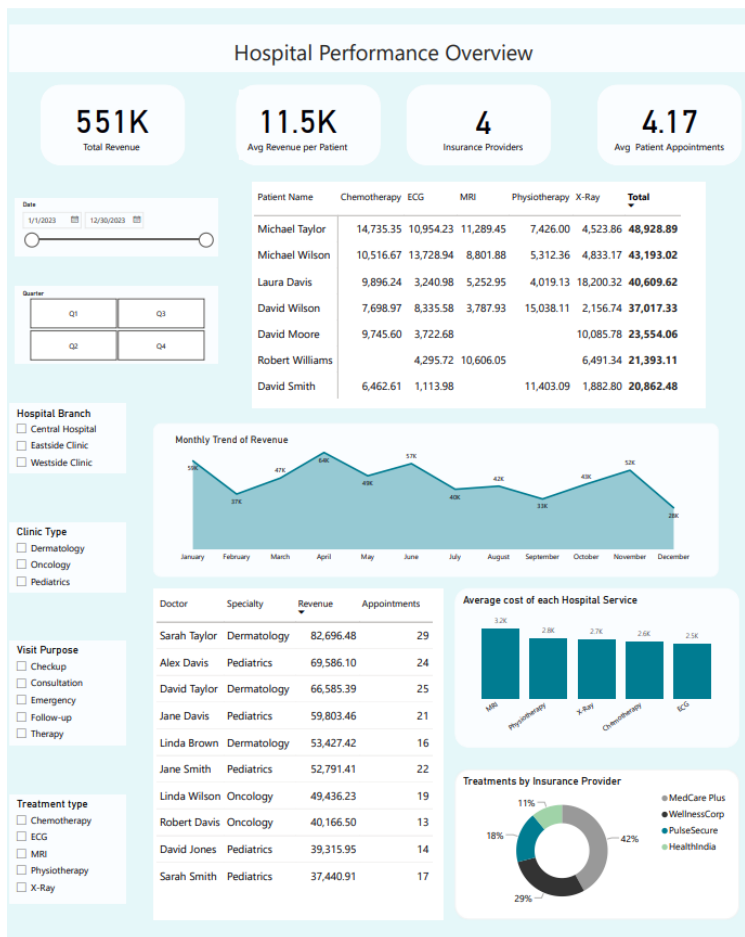
This page visualizes patient activity, no-show trends, and departmental efficiency to help identify operational bottlenecks and improve clinical workflow.

## Page 2: Financial & Strategic View

This page answers: “Where is our revenue coming from, and which services, doctors, and branches drive it?”

### Visuals:

- Total and average revenue cards
- Donut Chart: Insurance provider coverage (Medicare Plus covers 40%+)
- Column Chart: revenue by treatment (Chemotherapy is highest earner)
- Line Chart: monthly revenue fluctuations
- Table: Top 7 revenue-contributing patients and services used
- Table: Doctor specialization, appointments, patient volume, total revenue



### Financial Dashboard Page

This view presents financial performance by branch, doctor, and treatment type, enabling strategic revenue planning and resource allocation.

# Insights

This analysis uncovered key trends across clinical operations and financial performance:

**High No-Show Rates in Critical Services:** Departments like **X-ray (66%)** and **Emergency (~60%)** report the highest no-show rates. These are not only operational inefficiencies but potential risks for delayed care.

**Patient Care Fragmentation:** Several patients saw **3+ doctors**, and one consulted with **8 specialists**. This signals either complex health needs or disjointed care coordination both of which require process improvement and better patient tracking systems.

**Midweek Service Overload:** The **busiest appointment days are Tuesdays and Wednesdays**, creating potential bottlenecks in scheduling, staff availability, and service delivery. Spreading appointments more evenly could ease staff workload and enhance patient experience.

**Insurance Dependence:** All patients are insured, with **MedCare Plus alone covering 42%** of treatments. This dominance makes it a critical financial partner.

**Revenue Performance Gaps:** While **Chemotherapy** generates the most revenue, treatments like **MRI and Physiotherapy** are among the most expensive. This raises questions about pricing alignment, margins, and accessibility.

**Branch-Level Opportunity:** The **Central branch delivers 41.6% of total revenue**, yet lacks oncology; revealing a strong case for expanding high value services like chemotherapy to that location.

**Irregular Monthly Revenue:** Revenue shows volatile month over month trends. This may relate to seasonal patient flow, billing delays, or lack of proactive appointment targeting.

# Final Recommendations

**Reduce No-Shows in Key Departments:** Investigate root causes of absenteeism in X-ray and Emergency services. A mix of SMS/Email reminders, pre-visit calls, and workflow audits could boost attendance rates and operational throughput.

**Strengthen Care Coordination:** Patients seeing 5+ doctors may benefit from care navigation support or dedicated case managers. This will reduce redundancy, improve outcomes, and enhance patient satisfaction.

**Redistribute Appointment Load:** Encourage off-peak bookings through digital nudges or flexible hours to reduce the midweek bottleneck. This improves staff scheduling and balances demand across the week.

**Expand Oncology in High Performing Branches:** The Central branch's strong financial return without oncology indicates prime potential for expansion. A feasibility study on service rollout is recommended.

**Leverage Insurance Partnerships:** Collaborate with MedCare Plus and other insurers on bundled payments or preventive care incentives, as over 40% of patients are concentrated under one payer.

**Forecast & Stabilize Revenue Cycles:** Address month-to-month revenue swings by creating predictive models for appointment load and implementing consistent billing practices.

## **Conclusion**

This case study demonstrates how healthcare analytics can turn operational complexity into actionable insights. By connecting data across patient experiences, provider performance, and financial results, we can uncover patterns that inform both daily decisions and long term strategy.

From understanding what drives patient engagement to identifying the doctors and services fueling revenue, the analysis offers a full picture view of hospital performance, positioning leadership to be proactive, data driven, and aligned with care outcomes.